

Machine Learning 1

Nicole Manuguid

10/21/2021

First up is clustering methods

Kmeans clustering

The function in base R to do Kmeans clustering is called 'kmeans()'

First make up some data where we know what the answer should be:


rnorm gives you 30 values around $-/+3$

```
tmp <- c(rnorm(30,-3), rnorm(30, 3))
x <- cbind( x=tmp, y=rev(tmp))
x
```

```
##           x           y
## [1,] -3.1412044  3.0578624
## [2,] -3.4865367  4.0159913
## [3,] -2.5297047  3.5739667
## [4,] -3.7643103  0.8965555
## [5,] -2.5936659  2.7252673
## [6,] -3.2262401  3.4080059
## [7,] -4.2333415  1.9504400
## [8,] -1.9626672  0.9767131
## [9,] -2.4623646  3.1637550
## [10,] -2.6691961  2.2844080
## [11,] -2.5729675  3.3828995
## [12,] -2.4710376  3.5483131
## [13,] -2.4272385  3.7001909
## [14,] -2.4757287  2.5924046
## [15,] -2.7673957  1.8946373
## [16,] -3.2898923  2.7222560
## [17,] -2.0265646  4.2577452
## [18,] -3.6287904  2.4584898
## [19,] -3.4097228  4.4723335
## [20,] -2.1343836  3.3215864
## [21,] -1.5268920  3.5648982
## [22,] -3.9178263  3.8238795
## [23,] -1.8314844  4.0556140
## [24,] -2.0368922  3.2741820
## [25,] -3.5831635  4.2090731
## [26,] -2.8534937  3.0382699
```

```
## [27,] -2.8858429  2.9630653
## [28,] -3.3296444  3.7055410
## [29,] -2.7888602  2.1265879
## [30,] -2.3630077  4.8473191
## [31,]  4.8473191 -2.3630077
## [32,]  2.1265879 -2.7888602
## [33,]  3.7055410 -3.3296444
## [34,]  2.9630653 -2.8858429
## [35,]  3.0382699 -2.8534937
## [36,]  4.2090731 -3.5831635
## [37,]  3.2741820 -2.0368922
## [38,]  4.0556140 -1.8314844
## [39,]  3.8238795 -3.9178263
## [40,]  3.5648982 -1.5268920
## [41,]  3.3215864 -2.1343836
## [42,]  4.4723335 -3.4097228
## [43,]  2.4584898 -3.6287904
## [44,]  4.2577452 -2.0265646
## [45,]  2.7222560 -3.2898923
## [46,]  1.8946373 -2.7673957
## [47,]  2.5924046 -2.4757287
## [48,]  3.7001909 -2.4272385
## [49,]  3.5483131 -2.4710376
## [50,]  3.3828995 -2.5729675
## [51,]  2.2844080 -2.6691961
## [52,]  3.1637550 -2.4623646
## [53,]  0.9767131 -1.9626672
## [54,]  1.9504400 -4.2333415
## [55,]  3.4080059 -3.2262401
## [56,]  2.7252673 -2.5936659
## [57,]  0.8965555 -3.7643103
## [58,]  3.5739667 -2.5297047
## [59,]  4.0159913 -3.4865367
## [60,]  3.0578624 -3.1412044
```

```
plot(x)
```



Q. Can we use `kmeans()` to cluster this data setting `k` to 2 and `nstart` to 20?

```
km <- kmeans(x, centers = 2, nstart = 20)
km
```

[illegible]

Q. How many points are in each cluster?

```
km$size
```

```
## [1] 30 30
```

Q. What 'component' of your result object details cluster assignment/membership?

```
km$clusters
```

```
## NULL
```


Q. What 'component' of your result object details cluster center?

```
km$centers
```

```
##           x           y
## 1 -2.813002  3.133742
## 2  3.133742 -2.813002
```

Q. Plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
plot(x, col=km$cluster)
points(km$centers, col = "blue", pch = 15, cex = 2)
```




```
tmp <- c(rnorm(30,-3), rnorm(30, 3))
x <- cbind( x=tmp, y=rev(tmp))
x
```

```
##           x           y
## [1,] -2.6640839  2.6059780
## [2,] -2.9390800  2.0082212
## [3,] -1.5623976  5.5306240
## [4,] -3.7372025  3.4258740
## [5,] -3.3626717  3.6969259
## [6,] -2.7239710  3.4631122
## [7,] -3.7630841  2.5250689
## [8,] -3.6806079  3.6495002
## [9,] -3.0315489  2.8454496
## [10,] -0.7395924  3.6083428
## [11,] -3.2493730  1.6109370
## [12,] -4.9111490  2.4671732
## [13,] -3.9006485  2.9763371
## [14,] -2.2065966  2.8901086
## [15,] -2.9231745  4.3602302
## [16,] -4.3253700  2.5016956
## [17,] -4.5275763  1.8744552
## [18,] -2.7155145  3.1375565
## [19,] -3.8382030  3.8144844
## [20,] -4.2992578  3.3167316
## [21,] -2.0352245  2.1241205
## [22,] -2.5528483  3.0789149
## [23,] -1.4907307  3.4723631
## [24,] -3.7912991  4.3452170
## [25,] -3.2508148  4.1502938
## [26,] -2.2925818  3.5603762
## [27,] -3.3067187  2.6057904
## [28,] -3.9234862  1.8019638
## [29,] -3.1616756  1.7532017
## [30,] -2.5367984  2.9396486
## [31,]  2.9396486 -2.5367984
## [32,]  1.7532017 -3.1616756
## [33,]  1.8019638 -3.9234862
## [34,]  2.6057904 -3.3067187
## [35,]  3.5603762 -2.2925818
## [36,]  4.1502938 -3.2508148
## [37,]  4.3452170 -3.7912991
## [38,]  3.4723631 -1.4907307
## [39,]  3.0789149 -2.5528483
## [40,]  2.1241205 -2.0352245
## [41,]  3.3167316 -4.2992578
## [42,]  3.8144844 -3.8382030
## [43,]  3.1375565 -2.7155145
## [44,]  1.8744552 -4.5275763
## [45,]  2.5016956 -4.3253700
## [46,]  4.3602302 -2.9231745
## [47,]  2.8901086 -2.2065966
## [48,]  2.9763371 -3.9006485
```

```
## [49,] 2.4671732 -4.9111490
## [50,] 1.6109370 -3.2493730
## [51,] 3.6083428 -0.7395924
## [52,] 2.8454496 -3.0315489
## [53,] 3.6495002 -3.6806079
## [54,] 2.5250689 -3.7630841
## [55,] 3.4631122 -2.7239710
## [56,] 3.6969259 -3.3626717
## [57,] 3.4258740 -3.7372025
## [58,] 5.5306240 -1.5623976
## [59,] 2.0082212 -2.9390800
## [60,] 2.6059780 -2.6640839
```

```
plot(x)
```



Hierararchical Clustering

A big limitation with kmeans is that we have to tell it K (the number of clusters we want)

Analyze this same data with hclust()


Demonstrate the use of dist(), hclust(), plot (), and cutree() functions to do clustering, Generate dendrograms and return cluster assingment membership vector...

```
hc <- hclust(dist(x))
hc
```

```
##  
## Call:  
## hclust(d = dist(x))  
##  
## Cluster method   : complete  
## Distance         : euclidean  
## Number of objects: 60
```

There is a plot method for hclust result objects. Let's see it

```
plot(hc)
```



To get our cluster membership vector we have to do a wee bit more work. We have to “cut” the tree where we think it makes sense. For this we use the ‘`cutree()`’ function

```
cutree(hc, h = 6)
```


```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2  
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

You can also call 'cutree()' setting k=the number of grps/clusters you want.

```
grps <- cutree(hc, k = 2)
```

Make our results plot

```
plot(x, col=grps)
```



#Principal Component Analysis(PCA) of UK Food Data

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
## [1] 17 5
```

Checking our data

```
View(x)
```

fix row names; this should be 17 x 4 but we have 17 x5 dimensions


```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

```
##           England Wales Scotland N.Ireland
## Cheese      105   103     103      66
## Carcass_meat 245   227     242     267
## Other_meat   685   803     750     586
## Fish         147   160     122      93
## Fats_and_oils 193   235     184     209
## Sugars       156   175     147     139
```

recheck dimensions

```
dim(x)
```

```
## [1] 17  4
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?


be careful with previous function. can use this function below to get the same results

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
head(x)
```

```
##           England Wales Scotland N.Ireland
## Cheese      105   103     103      66
## Carcass_meat 245   227     242     267
## Other_meat   685   803     750     586
## Fish         147   160     122      93
## Fats_and_oils 193   235     184     209
## Sugars       156   175     147     139
```


Spotting major differences and trends

```
barplot(as.matrix(x), beside = T, col=rainbow(nrow(x)))
```




Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

```
barplot(as.matrix(x), beside = FALSE, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch = 16)
```



If a point lies in the diagonal between two countries it means they share similarities within the data.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

For N.Ireland, the blue point is not following the general shape of the diagonal and is most different compared to other countries in the UK.

#PCA to the rescue

the main function is in base R

the main function in base R for PCA is 'prcomp()' this wants the transpose of our data

t(x)

```
##      Cheese Carcass_meat Other_meat Fish Fats_and_oils Sugars
## England      105         245       685  147           193   156
## Wales        103         227       803  160           235   175
## Scotland     103         242       750  122           184   147
## N.Ireland     66         267       586   93           209   139
##
##      Fresh_potatoes Fresh_Veg Other_Veg Processed_potatoes
## England           720       253       488             198
## Wales             874       265       570             203
## Scotland          566       171       418             220
## N.Ireland         1033       143       355             187
##
##      Processed_Veg Fresh_fruit Cereals Beverages Soft_drinks
```

```
## England      360      1102      1472      57      1374
## Wales        365      1137      1582      73      1256
## Scotland     337      957      1462      53      1572
## N.Ireland    334      674      1494      47      1506
##      Alcoholic_drinks  Confectionery
## England      375      54
## Wales        475      64
## Scotland     458      62
## N.Ireland    135      41
```


```
pca <- prcomp(t(x))
summary(pca)
```

```
## Importance of components:
```

```
##      PC1      PC2      PC3      PC4
## Standard deviation 324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance 0.6744 0.2905 0.03503 0.000e+00
## Cumulative Proportion 0.6744 0.9650 1.00000 1.000e+00
```


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim =c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim = c(-270, 500))
text(pca$x[,1], pca$x[,2], colnames(x), col = c("yellow", "red", "blue", "green"))
```



calculate how much variation in the original data each PC accounts for


```
v <- round(pca$sdev^2/sum(pca$sdev^2)*100)
v
```

```
## [1] 67 29 4 0
```


```
z <- summary(pca)
z$importance
```

```
##              PC1      PC2      PC3      PC4
## Standard deviation 324.15019 212.74780 73.87622 4.188568e-14
## Proportion of Variance 0.67444 0.29052 0.03503 0.000000e+00
## Cumulative Proportion 0.67444 0.96497 1.00000 1.000000e+00
```

```
barplot(v, xlab = "Principal Component", ylab = "Percent Variation")
```




```
par(mar=c(10, 3, 0.35, 0))  
barplot (pca$rotation[,1], las=2)
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?


```
par(mar=c(10, 3, 0.35, 0))
barplot (pca$rotation[,2], las=2)
```

The two main food groups are fresh potatoes and soft drinks. PC2 shows us how much variation there is in food in N.Ireland compared to the rest of the countries in the UK.

#biplots

```
biplot(pca)
```



```
#pca of r-seq data
```

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

##	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
## gene1	439	458	408	429	420	90	88	86	90	93
## gene2	219	200	204	210	187	427	423	434	433	426
## gene3	1006	989	1030	1017	973	252	237	238	226	210
## gene4	783	792	829	856	760	849	856	835	885	894
## gene5	181	249	204	244	225	277	305	272	270	279
## gene6	460	502	491	491	493	612	594	577	618	638


Q10: How many genes and samples are in this data set?

```
dim(rna.data)
```

```
## [1] 100 10
```

There are 100 genes and 10 samples

```
pca <- prcomp(t(rna.data), scale = TRUE)
plot(pca$x[,1], pca$x[,2], xlab = "PC1", ylab = "PC2")
```



```
summary(pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  9.6237 1.5198 1.05787 1.05203 0.88062 0.82545 0.80111
## Proportion of Variance 0.9262 0.0231 0.01119 0.01107 0.00775 0.00681 0.00642
## Cumulative Proportion 0.9262 0.9493 0.96045 0.97152 0.97928 0.98609 0.99251
##              PC8      PC9      PC10
## Standard deviation  0.62065 0.60342 3.348e-15
## Proportion of Variance 0.00385 0.00364 0.000e+00
## Cumulative Proportion 0.99636 1.00000 1.000e+00
```

```
plot (pca, main = "Quick scree plot")
```

Quick scree plot



Variance captured per PC

```
pca.var <- pca$sdev^2
```

Percent variance is more informative to look at

```
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
## [1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```

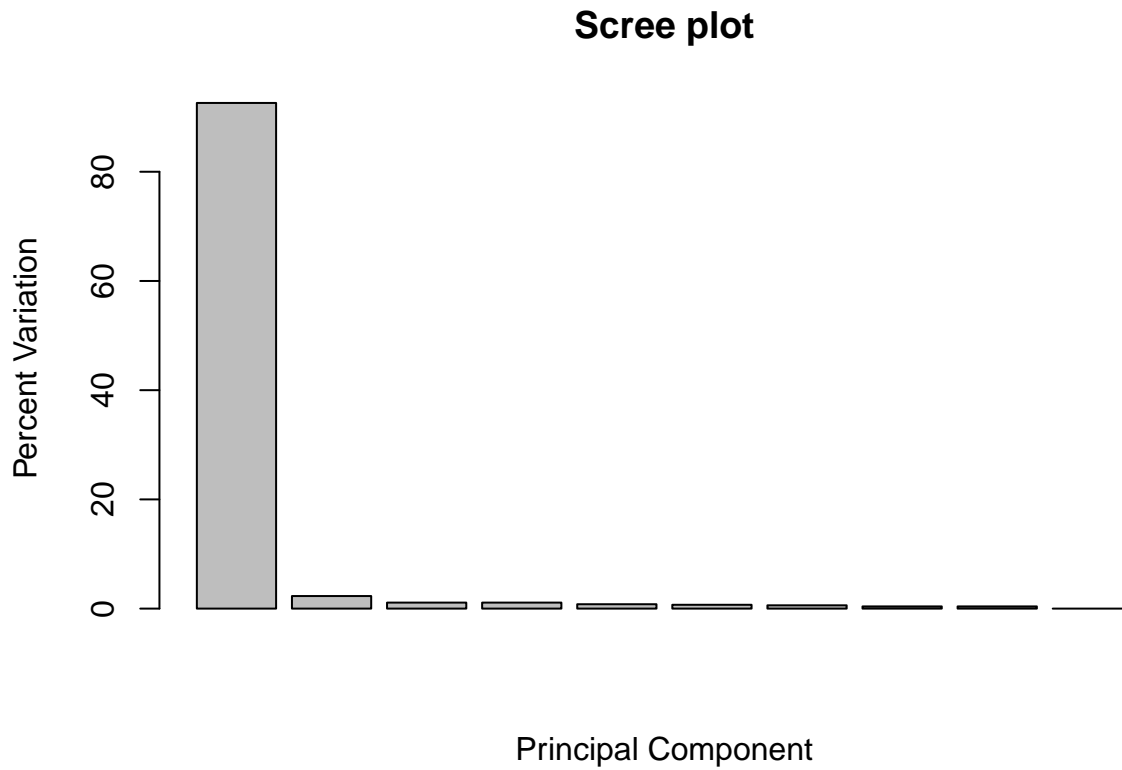
generate scree-plot

```
barplot(pca.var.per, main="Scree plot", name.arg = paste0 ("PC", 1:10), xlab = "Principal Component", ylab = "Percent Variance Captured")
```

```
## Warning in plot.window(xlim, ylim, log = log, ...): "name.arg" is not a
## graphical parameter
```

```
## Warning in title(main = main, sub = sub, xlab = xlab, ylab = ylab, ...):
## "name.arg" is not a graphical parameter
```


```
## Warning in axis(if (horiz) 1 else 2, cex.axis = cex.axis, ...): "name.arg" is
## not a graphical parameter
```



A vector of colors for wt and ko samples

```
colvec <- colnames(rna.data)
colvec[grep ("wt", colvec)] <- "red"
colvec[grep ("ko", colvec)] <- "blue"


plot(pca$x[,1], pca$x[,2], col = colvec, pch = 16, xlab = paste0("PC1 (", pca.var.per[1], "%"), ylab = "PC2 (",
text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```



text


```
## function (x, ...)\n## UseMethod("text")\n## <bytecode: 0x000000000f9785a8>\n## <environment: namespace:graphics>
```

```
library(ggplot2)\ndf <- as.data.frame(pca$x)\n\n#make first plot\n\n ggplot(df) + aes(PC1, PC2) + geom_point()
```



```
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data), 1,2)

p <- ggplot(df) + aes(PC1, PC2, label = samples, col = condition) + geom_label(show.legend = FALSE)
p
```




polish plot

```
p + labs(title = "PCA of RNASeq Data", subtitle = "PC1 clearly separates wild-type from knock-out samples")
```


PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



BIMM143 example data